Deep Generative Models

Lecture 11

Roman Isachenko

Moscow Institute of Physics and Technology
Yandex School of Data Analysis

2025, Autumn

Recap of Previous Lecture

$$\frac{d\mathbf{x}(t)}{dt} = \mathbf{f}_\theta(\mathbf{x}(t), t); \quad \text{with initial condition } \mathbf{x}(t_0) = \mathbf{x}_0$$

Theorem (Continuity Equation)

If $\mathbf{f}$ is uniformly Lipschitz continuous in $\mathbf{x}$ and continuous in t, then

$$\frac{d \log p_t(\mathbf{x}(t))}{dt} = -\text{tr}\left(\frac{\partial \mathbf{f}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)}\right)$$

Solution of the Continuity Equation

$$\log p_1(\mathbf{x}(1)) = \log p_0(\mathbf{x}(0)) - \int_0^1 \text{tr}\left(\frac{\partial \mathbf{f}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)}\right) dt.$$

▶ This solution gives us the density along the trajectory (not the total probability path).

▶ However, it's difficult to efficiently estimate the last term.

Chen R. T. Q. et al. Neural Ordinary Differential Equations, 2018

Recap of Previous Lecture

SDE Basics

Let's define a stochastic process $\mathbf{x}(t)$ with initial condition $\mathbf{x}(0) \sim p_0(\mathbf{x})$:

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w},$$

where $\mathbf{w}(t)$ is the standard Wiener process (Brownian motion):

$$\mathbf{w}(t) - \mathbf{w}(s) \sim \mathcal{N}(0, (t-s)\mathbf{I}), \quad d\mathbf{w} = \boldsymbol{\epsilon} \cdot \sqrt{dt}, \text{ where } \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I}).$$

Discretization of SDE (Euler Method) - SDESolve

$$\mathbf{x}(t + dt) = \mathbf{x}(t) + \mathbf{f}(\mathbf{x}(t), t) \cdot dt + g(t) \cdot \boldsymbol{\epsilon} \cdot \sqrt{dt}$$

▶ At each time t, we have the density $p_t(\mathbf{x}) = p(\mathbf{x}, t)$.

▶ $p : \mathbb{R}^m \times [0, 1] \to \mathbb{R}_+$ is a **probability path** between $p_0(\mathbf{x})$ and $p_1(\mathbf{x})$.

Recap of Previous Lecture

Theorem (Kolmogorov-Fokker-Planck)

The evolution of the distribution $p_t(\mathbf{x})$ is given by:

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x})\right) + \frac{1}{2}g^2(t)\Delta_{\mathbf{x}}p_t(\mathbf{x})$$

Langevin SDE (Special Case)

$$d\mathbf{x} = \frac{1}{2}\frac{\partial}{\partial \mathbf{x}}\log p_t(\mathbf{x})dt + 1 \cdot d\mathbf{w}$$

The density $p_\theta(\mathbf{x})$ is a **stationary** distribution for the SDE.

Langevin Dynamics

Samples from the following dynamics will come from $p_\theta(\mathbf{x})$ under mild regularity conditions for a small enough η:

$$\mathbf{x}_{t+1} = \mathbf{x}_t + \frac{\eta}{2}\nabla_{\mathbf{x}_t}\log p_\theta(\mathbf{x}_t) + \sqrt{\eta} \cdot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I}).$$

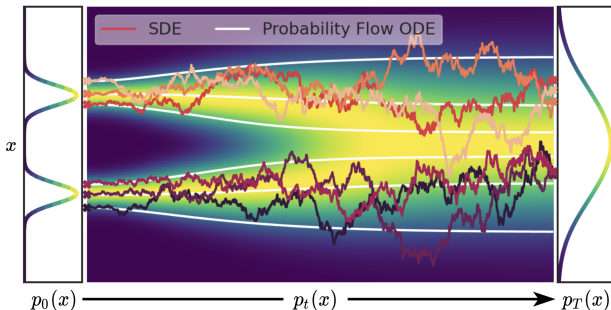*Welling M. Bayesian Learning via Stochastic Gradient Langevin Dynamics, 2011*

Recap of Previous Lecture

$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ (SDE with the probability path $p_t(\mathbf{x})$)

Probability Flow ODE

There exists an ODE with the identical probability path $p_t(\mathbf{x})$ of the form:

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt$$

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Recap of Previous Lecture

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt, \quad \mathbf{x}(t + dt) = \mathbf{x}(t) + \mathbf{f}(\mathbf{x}, t)dt$$

Reverse ODE
Let $\tau = 1 - t$ ($d\tau = -dt$):

$$d\mathbf{x} = -\mathbf{f}(\mathbf{x}, 1 - \tau)d\tau$$

Reverse SDE
There's a reverse SDE for $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ in the following form:

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - g^2(t)\frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt + g(t)d\mathbf{w}, \quad dt < 0$$
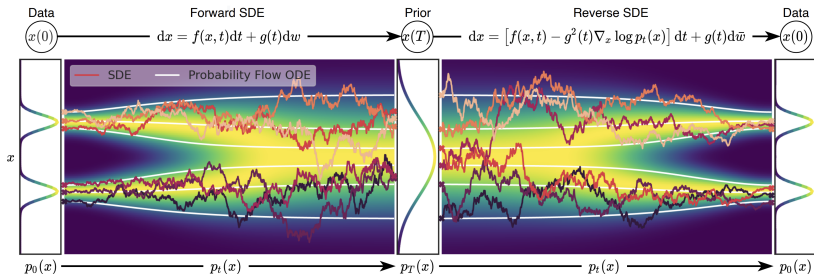
Sketch of the Proof

▶ Convert the initial SDE to the probability flow ODE.

▶ Reverse the probability flow ODE.

▶ Convert the reverse probability flow ODE to the reverse SDE.

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Recap of Previous Lecture

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w} \qquad \text{(SDE)}$$

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\frac{\partial}{\partial \mathbf{x}}\log p_t(\mathbf{x}) \right) dt \quad \text{(probability flow ODE)}$$

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - g^2(t)\frac{\partial}{\partial \mathbf{x}}\log p_t(\mathbf{x}) \right) dt + g(t)d\mathbf{w} \quad \text{(reverse SDE)}$$

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Outline

Outline

Score Matching SDE

Denoising Score Matching

$$\mathbf{x}_t = \mathbf{x} + \sigma_t \cdot \boldsymbol{\epsilon}_t, \qquad q(\mathbf{x}_t|\mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma_t^2 \cdot \mathbf{I})$$

$$\mathbf{x}_{t-1} = \mathbf{x} + \sigma_{t-1} \cdot \boldsymbol{\epsilon}_{t-1}, \qquad q(\mathbf{x}_{t-1}|\mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma_{t-1}^2 \cdot \mathbf{I})$$

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Score Matching SDE

Denoising Score Matching

$$\mathbf{x}_t = \mathbf{x} + \sigma_t \cdot \boldsymbol{\epsilon}_t, \qquad\qquad q(\mathbf{x}_t|\mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma_t^2 \cdot \mathbf{I})$$

$$\mathbf{x}_{t-1} = \mathbf{x} + \sigma_{t-1} \cdot \boldsymbol{\epsilon}_{t-1}, \qquad q(\mathbf{x}_{t-1}|\mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma_{t-1}^2 \cdot \mathbf{I})$$

$$\mathbf{x}_t = \mathbf{x}_{t-1} + \sqrt{\sigma_t^2 - \sigma_{t-1}^2} \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_{t-1}, (\sigma_t^2 - \sigma_{t-1}^2) \cdot \mathbf{I})$$

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Score Matching SDE

Denoising Score Matching

$$\mathbf{x}_t = \mathbf{x} + \sigma_t \cdot \boldsymbol{\epsilon}_t, \qquad q(\mathbf{x}_t|\mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma_t^2 \cdot \mathbf{I})$$

$$\mathbf{x}_{t-1} = \mathbf{x} + \sigma_{t-1} \cdot \boldsymbol{\epsilon}_{t-1}, \qquad q(\mathbf{x}_{t-1}|\mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma_{t-1}^2 \cdot \mathbf{I})$$

$$\mathbf{x}_t = \mathbf{x}_{t-1} + \sqrt{\sigma_t^2 - \sigma_{t-1}^2} \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_{t-1}, (\sigma_t^2 - \sigma_{t-1}^2) \cdot \mathbf{I})$$

Let's transform this Markov chain into the continuous stochastic process $\mathbf{x}(t)$ by letting $T \to \infty$:

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\sigma^2(t) - \sigma^2(t - dt)} \cdot \boldsymbol{\epsilon}$$

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Score Matching SDE

Denoising Score Matching

$$\mathbf{x}_t = \mathbf{x} + \sigma_t \cdot \boldsymbol{\epsilon}_t, \qquad\qquad q(\mathbf{x}_t|\mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma_t^2 \cdot \mathbf{I})$$

$$\mathbf{x}_{t-1} = \mathbf{x} + \sigma_{t-1} \cdot \boldsymbol{\epsilon}_{t-1}, \qquad q(\mathbf{x}_{t-1}|\mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma_{t-1}^2 \cdot \mathbf{I})$$

$$\mathbf{x}_t = \mathbf{x}_{t-1} + \sqrt{\sigma_t^2 - \sigma_{t-1}^2} \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_{t-1}, (\sigma_t^2 - \sigma_{t-1}^2) \cdot \mathbf{I})$$

Let's transform this Markov chain into the continuous stochastic process $\mathbf{x}(t)$ by letting $T \to \infty$:

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\sigma^2(t) - \sigma^2(t - dt)} \cdot \boldsymbol{\epsilon}$$

$$= \mathbf{x}(t - dt) + \sqrt{\frac{\sigma^2(t) - \sigma^2(t - dt)}{dt} dt} \cdot \boldsymbol{\epsilon}$$

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Score Matching SDE

Denoising Score Matching

$$\mathbf{x}_t = \mathbf{x} + \sigma_t \cdot \boldsymbol{\epsilon}_t, \qquad q(\mathbf{x}_t|\mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma_t^2 \cdot \mathbf{I})$$

$$\mathbf{x}_{t-1} = \mathbf{x} + \sigma_{t-1} \cdot \boldsymbol{\epsilon}_{t-1}, \qquad q(\mathbf{x}_{t-1}|\mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma_{t-1}^2 \cdot \mathbf{I})$$

$$\mathbf{x}_t = \mathbf{x}_{t-1} + \sqrt{\sigma_t^2 - \sigma_{t-1}^2} \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_{t-1}, (\sigma_t^2 - \sigma_{t-1}^2) \cdot \mathbf{I})$$

Let's transform this Markov chain into the continuous stochastic process $\mathbf{x}(t)$ by letting $T \to \infty$:

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\sigma^2(t) - \sigma^2(t - dt)} \cdot \boldsymbol{\epsilon}$$

$$= \mathbf{x}(t - dt) + \sqrt{\frac{\sigma^2(t) - \sigma^2(t - dt)}{dt} dt} \cdot \boldsymbol{\epsilon}$$

$$= \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Score Matching SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Score Matching SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

Variance Exploding SDE

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

$\sigma(t)$ is a monotonically increasing function.

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Score Matching SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

Variance Exploding SDE

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

$\sigma(t)$ is a monotonically increasing function.

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}, \quad \mathbf{f}(\mathbf{x}, t) = 0, \quad g(t) = \sqrt{\frac{d[\sigma^2(t)]}{dt}}$$

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Score Matching SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

Variance Exploding SDE

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

$\sigma(t)$ is a monotonically increasing function.

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}, \quad \mathbf{f}(\mathbf{x}, t) = 0, \quad g(t) = \sqrt{\frac{d[\sigma^2(t)]}{dt}}$$

$$d\mathbf{x} = \left(-\frac{1}{2}\frac{d[\sigma^2(t)]}{dt}\frac{\partial}{\partial \mathbf{x}}\log p_t(\mathbf{x}) \right) dt \qquad \text{(probability flow ODE)}$$

$$d\mathbf{x} = \left(-\frac{d[\sigma^2(t)]}{dt}\frac{\partial}{\partial \mathbf{x}}\log p_t(\mathbf{x}) \right) dt + \sqrt{\frac{d[\sigma^2(t)]}{dt}}d\mathbf{w} \quad \text{(reverse SDE)}$$

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Diffusion SDE

Denoising Diffusion

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1} + \sqrt{\beta_t} \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1}, \beta_t \cdot \mathbf{I})$$

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Diffusion SDE

Denoising Diffusion

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1} + \sqrt{\beta_t} \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1}, \beta_t \cdot \mathbf{I})$$

Let's turn this Markov chain into a continuous stochastic process by letting $T \to \infty$ and setting $\beta_t = \beta(\frac{t}{T}) \cdot \frac{1}{T}$ (where $dt = \frac{1}{T}$):

$$\mathbf{x}(t) = \sqrt{1 - \beta(t)dt} \cdot \mathbf{x}(t - dt) + \sqrt{\beta(t)dt} \cdot \boldsymbol{\epsilon}$$

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Diffusion SDE

Denoising Diffusion

$$\mathbf{x}_t = \sqrt{1-\beta_t}\cdot\mathbf{x}_{t-1} + \sqrt{\beta_t}\cdot\boldsymbol{\epsilon}, \quad q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1-\beta_t}\cdot\mathbf{x}_{t-1}, \beta_t\cdot\mathbf{I})$$

Let's turn this Markov chain into a continuous stochastic process by letting $T \to \infty$ and setting $\beta_t = \beta(\frac{t}{T}) \cdot \frac{1}{T}$ (where $dt = \frac{1}{T}$):

$$\mathbf{x}(t) = \sqrt{1 - \beta(t)dt} \cdot \mathbf{x}(t - dt) + \sqrt{\beta(t)dt} \cdot \boldsymbol{\epsilon} \approx$$
$$\approx (1 - \frac{1}{2}\beta(t)dt) \cdot \mathbf{x}(t - dt) + \sqrt{\beta(t)dt} \cdot \boldsymbol{\epsilon}$$

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Diffusion SDE

Denoising Diffusion

$$\mathbf{x}_t = \sqrt{1-\beta_t} \cdot \mathbf{x}_{t-1} + \sqrt{\beta_t} \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1-\beta_t} \cdot \mathbf{x}_{t-1}, \beta_t \cdot \mathbf{I})$$

Let's turn this Markov chain into a continuous stochastic process by letting $T \to \infty$ and setting $\beta_t = \beta(\frac{t}{T}) \cdot \frac{1}{T}$ (where $dt = \frac{1}{T}$):

$$\mathbf{x}(t) = \sqrt{1-\beta(t)dt} \cdot \mathbf{x}(t-dt) + \sqrt{\beta(t)dt} \cdot \boldsymbol{\epsilon} \approx$$
$$\approx (1 - \frac{1}{2}\beta(t)dt) \cdot \mathbf{x}(t-dt) + \sqrt{\beta(t)dt} \cdot \boldsymbol{\epsilon} =$$
$$= \mathbf{x}(t-dt) - \frac{1}{2}\beta(t)\mathbf{x}(t-dt)dt + \sqrt{\beta(t)} \cdot d\mathbf{w}$$

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Diffusion SDE

Denoising Diffusion

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1} + \sqrt{\beta_t} \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1}, \beta_t \cdot \mathbf{I})$$

Let's turn this Markov chain into a continuous stochastic process by letting $T \to \infty$ and setting $\beta_t = \beta(\frac{t}{T}) \cdot \frac{1}{T}$ (where $dt = \frac{1}{T}$):

$$\mathbf{x}(t) = \sqrt{1 - \beta(t)dt} \cdot \mathbf{x}(t - dt) + \sqrt{\beta(t)dt} \cdot \boldsymbol{\epsilon} \approx$$
$$\approx (1 - \frac{1}{2}\beta(t)dt) \cdot \mathbf{x}(t - dt) + \sqrt{\beta(t)dt} \cdot \boldsymbol{\epsilon} =$$
$$= \mathbf{x}(t - dt) - \frac{1}{2}\beta(t)\mathbf{x}(t - dt)dt + \sqrt{\beta(t)} \cdot d\mathbf{w}$$

Variance Preserving SDE

$$d\mathbf{x} = -\frac{1}{2}\beta(t)\mathbf{x}(t)dt + \sqrt{\beta(t)} \cdot d\mathbf{w}$$

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Diffusion SDE

Variance Preserving SDE

$$dx = -\frac{1}{2}\beta(t)x(t)dt + \sqrt{\beta(t)} \cdot dw$$

$$f(x, t) = -\frac{1}{2}\beta(t)x(t), \quad g(t) = \sqrt{\beta(t)}$$

Variance is preserved as long as $x(0)$ has unit variance.

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Diffusion SDE

Variance Preserving SDE

$$d\mathbf{x} = -\frac{1}{2}\beta(t)\mathbf{x}(t)dt + \sqrt{\beta(t)} \cdot d\mathbf{w}$$

$$\mathbf{f}(\mathbf{x}, t) = -\frac{1}{2}\beta(t)\mathbf{x}(t), \quad g(t) = \sqrt{\beta(t)}$$

Variance is preserved as long as $\mathbf{x}(0)$ has unit variance.

$$d\mathbf{x} = \left(-\frac{1}{2}\beta(t)\mathbf{x}(t) - \frac{1}{2}\beta(t)\frac{\partial}{\partial\mathbf{x}}\log p_t(\mathbf{x}) \right) dt \qquad \text{(probability flow ODE)}$$

$$d\mathbf{x} = \left(-\frac{1}{2}\beta(t)\mathbf{x}(t) - \beta(t)\frac{\partial}{\partial\mathbf{x}}\log p_t(\mathbf{x}) \right) dt + \sqrt{\beta(t)}d\mathbf{w} \quad \text{(reverse SDE)}$$

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Diffusion SDE

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

Lu C. et al. Dpm-solver: A fast ode solver for diffusion probabilistic model sampling in around 10 steps, 2022

Diffusion SDE

$$dx = f(x, t)dt + g(t)dw$$

Variance Exploding SDE (NCSN)

$$dx = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot dw$$

Variance Preserving SDE (DDPM)

$$dx = -\frac{1}{2}\beta(t)x(t)dt + \sqrt{\beta(t)} \cdot dw$$

Lu C. et al. Dpm-solver: A fast ode solver for diffusion probabilistic model sampling in around 10 steps, 2022

Diffusion SDE

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

Variance Exploding SDE (NCSN)

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

Variance Preserving SDE (DDPM)

$$d\mathbf{x} = -\frac{1}{2}\beta(t)\mathbf{x}(t)dt + \sqrt{\beta(t)} \cdot d\mathbf{w}$$

Efficient Solvers

▶ Converting SDEs to PF-ODEs yields more efficient inference.

▶ We can apply any `ODESolve` procedure to reduce the number of inference steps.

▶ In practice, this reduces the number of steps from 100–1000 to 20–50.

Lu C. et al. Dpm-solver: A fast ode solver for diffusion probabilistic model sampling in around 10 steps, 2022

Outline

Score-Based Generative Models Through SDEs

Discrete-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}_0)}\mathbb{E}_{t\sim U\{1,T\}}\mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)}\left\|\mathbf{s}_{\boldsymbol{\theta},t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t}\log q(\mathbf{x}_t|\mathbf{x}_0)\right\|_2^2$$

Is it possible to train score-based diffusion models in continuous time?

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Score-Based Generative Models Through SDEs

Discrete-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}_0)} \mathbb{E}_{t \sim U\{1,T\}} \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} \left\| \mathbf{s}_{\boldsymbol{\theta},t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}_0) \right\|_2^2$$

Is it possible to train score-based diffusion models in continuous time?

Continuous-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))} \mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))} \left\| \mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}(t), t) - \nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t)|\mathbf{x}(0)) \right\|_2^2$$

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Score-Based Generative Models Through SDEs

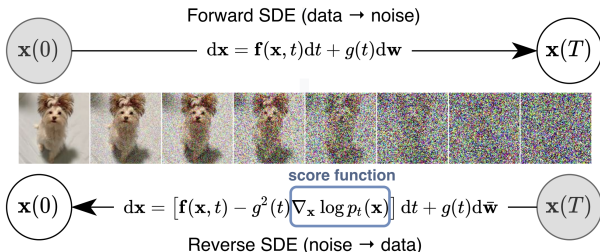## Discrete-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}_0)} \mathbb{E}_{t \sim U\{1,T\}} \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} \left\| \mathbf{s}_{\boldsymbol{\theta},t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}_0) \right\|_2^2$$

Is it possible to train score-based diffusion models in continuous time?

Continuous-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))} \mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))} \left\| \mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}(t),t) - \nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t)|\mathbf{x}(0)) \right\|_2^2$$

Forward SDE (data → noise)

$\mathbf{x}(0)$ ——— $d\mathbf{x} = \mathbf{f}(\mathbf{x},t)dt + g(t)d\mathbf{w}$ ——→ $\mathbf{x}(T)$

score function

$\mathbf{x}(0)$ ←—— $d\mathbf{x} = \left[\mathbf{f}(\mathbf{x},t) - g^2(t) \boxed{\nabla_{\mathbf{x}} \log p_t(\mathbf{x})} \right] dt + g(t)d\bar{\mathbf{w}}$ ——— $\mathbf{x}(T)$

Reverse SDE (noise → data)

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Score-Based Generative Models Through SDEs

Continuous-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))}\mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))}\big\|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}(t),t)-\nabla_{\mathbf{x}(t)}\log q(\mathbf{x}(t)|\mathbf{x}(0))\big\|_2^2$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

Continuous-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))}\mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))}\big\|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}(t),t)-\nabla_{\mathbf{x}(t)}\log q(\mathbf{x}(t)|\mathbf{x}(0))\big\|_2^2$$

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\Big(\boldsymbol{\mu}(\mathbf{x}(t),\mathbf{x}(0)), \boldsymbol{\sigma}^2(\mathbf{x}(t),\mathbf{x}(0))\cdot\mathbf{I}\Big)$$

$$\nabla_{\mathbf{x}(t)}\log q(\mathbf{x}(t)|\mathbf{x}(0)) = -\frac{1}{\boldsymbol{\sigma}}\odot(\mathbf{x}(t)-\boldsymbol{\mu})$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

Continuous-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))}\mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))}\big\|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}(t),t)-\nabla_{\mathbf{x}(t)}\log q(\mathbf{x}(t)|\mathbf{x}(0))\big\|_2^2$$

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\Big(\boldsymbol{\mu}(\mathbf{x}(t),\mathbf{x}(0)), \boldsymbol{\sigma}^2(\mathbf{x}(t),\mathbf{x}(0))\cdot\mathbf{I}\Big)$$

$$\nabla_{\mathbf{x}(t)}\log q(\mathbf{x}(t)|\mathbf{x}(0)) = -\frac{1}{\boldsymbol{\sigma}}\odot(\mathbf{x}(t)-\boldsymbol{\mu})$$

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}}\cdot d\mathbf{w} \quad \text{(Variance Exploding SDE)}$$

$$d\mathbf{x} = -\frac{1}{2}\beta(t)\mathbf{x}(t)dt + \sqrt{\beta(t)}\cdot d\mathbf{w} \quad \text{(Variance Preserving SDE)}$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

Continuous-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))}\mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))}\big\|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}(t),t)-\nabla_{\mathbf{x}(t)}\log q(\mathbf{x}(t)|\mathbf{x}(0))\big\|_2^2$$

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\Big(\boldsymbol{\mu}(\mathbf{x}(t),\mathbf{x}(0)), \boldsymbol{\sigma}^2(\mathbf{x}(t),\mathbf{x}(0))\cdot\mathbf{I}\Big)$$

$$\nabla_{\mathbf{x}(t)}\log q(\mathbf{x}(t)|\mathbf{x}(0)) = -\frac{1}{\boldsymbol{\sigma}}\odot(\mathbf{x}(t)-\boldsymbol{\mu})$$

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}}\cdot d\mathbf{w} \quad \text{(Variance Exploding SDE)}$$

$$d\mathbf{x} = -\frac{1}{2}\beta(t)\mathbf{x}(t)dt + \sqrt{\beta(t)}\cdot d\mathbf{w} \quad \text{(Variance Preserving SDE)}$$

Is it possible to explicitly derive $\boldsymbol{\mu}(\mathbf{x}(t),\mathbf{x}(0))$ and $\boldsymbol{\Sigma}(\mathbf{x}(t),\mathbf{x}(0))$ for VE-SDE and VP-SDE?

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\Big(\boldsymbol{\mu}(\mathbf{x}(t),\mathbf{x}(0)),\boldsymbol{\Sigma}(\mathbf{x}(t),\mathbf{x}(0))\Big)$$

Theorem

The moments of the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x},t)dt + g(t)d\mathbf{w}$ satisfy:

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(t),\mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}(t),t)|\mathbf{x}(0)\right]$$

$$\frac{d\boldsymbol{\Sigma}(\mathbf{x}(t),\mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}\cdot(\mathbf{x}(t)-\boldsymbol{\mu})^T + (\mathbf{x}(t)-\boldsymbol{\mu})\cdot\mathbf{f}^T|\mathbf{x}(0)\right] + g^2(t)\cdot\mathbf{I}$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\Big(\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0)), \boldsymbol{\Sigma}(\mathbf{x}(t), \mathbf{x}(0))\Big)$$

Theorem

The moments of the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ satisfy:

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}(t), t)|\mathbf{x}(0)\right]$$

$$\frac{d\boldsymbol{\Sigma}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f} \cdot (\mathbf{x}(t) - \boldsymbol{\mu})^T + (\mathbf{x}(t) - \boldsymbol{\mu}) \cdot \mathbf{f}^T|\mathbf{x}(0)\right] + g^2(t)\cdot\mathbf{I}$$

Proof

$$\mathbb{E}\left[d\mathbf{x}|\mathbf{x}(0)\right] = \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)dt|\mathbf{x}(0)\right] + \mathbb{E}\left[g(t)d\mathbf{w}|\mathbf{x}(0)\right]$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\Big(\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0)), \boldsymbol{\Sigma}(\mathbf{x}(t), \mathbf{x}(0))\Big)$$

Theorem

The moments of the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ satisfy:

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}(t), t)|\mathbf{x}(0)\right]$$

$$\frac{d\boldsymbol{\Sigma}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f} \cdot (\mathbf{x}(t) - \boldsymbol{\mu})^T + (\mathbf{x}(t) - \boldsymbol{\mu}) \cdot \mathbf{f}^T | \mathbf{x}(0)\right] + g^2(t) \cdot \mathbf{I}$$

Proof

$$\mathbb{E}\left[d\mathbf{x}|\mathbf{x}(0)\right] = \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)dt|\mathbf{x}(0)\right] + \mathbb{E}\left[g(t)d\mathbf{w}|\mathbf{x}(0)\right]$$

$$= \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right]dt + g(t)\mathbb{E}\left[d\mathbf{w}|\mathbf{x}(0)\right]$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\Big(\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0)), \boldsymbol{\Sigma}(\mathbf{x}(t), \mathbf{x}(0))\Big)$$

Theorem

The moments of the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ satisfy:

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}(t), t)|\mathbf{x}(0)\right]$$

$$\frac{d\boldsymbol{\Sigma}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f} \cdot (\mathbf{x}(t) - \boldsymbol{\mu})^T + (\mathbf{x}(t) - \boldsymbol{\mu}) \cdot \mathbf{f}^T|\mathbf{x}(0)\right] + g^2(t) \cdot \mathbf{I}$$

Proof

$$\begin{aligned}
\mathbb{E}\left[d\mathbf{x}|\mathbf{x}(0)\right] &= \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)dt|\mathbf{x}(0)\right] + \mathbb{E}\left[g(t)d\mathbf{w}|\mathbf{x}(0)\right] \\
&= \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right] dt + g(t)\mathbb{E}\left[d\mathbf{w}|\mathbf{x}(0)\right] \\
&= \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right] dt
\end{aligned}$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}(t), t)|\mathbf{x}(0)\right]$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}(t), t)|\mathbf{x}(0)\right]$$

Proof (Continued)

$$\mathbb{E}\left[d\mathbf{x}|\mathbf{x}(0)\right] = \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right]dt$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}(t), t)|\mathbf{x}(0)\right]$$

Proof (Continued)

$$\mathbb{E}\left[d\mathbf{x}|\mathbf{x}(0)\right] = \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right] dt$$

$$\frac{d\mathbb{E}\left[\mathbf{x}(t)|\mathbf{x}(0)\right]}{dt} = \frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right]$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}(t), t)|\mathbf{x}(0)\right]$$

Proof (Continued)

$$\mathbb{E}\left[d\mathbf{x}|\mathbf{x}(0)\right] = \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right] dt$$

$$\frac{d\mathbb{E}\left[\mathbf{x}(t)|\mathbf{x}(0)\right]}{dt} = \frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right]$$

Examples

$$\textbf{NCSN:} \quad \mathbf{f}(\mathbf{x}, t) = 0 \quad \Rightarrow \quad \boldsymbol{\mu} = \mathbf{x}(0)$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}(t), t)|\mathbf{x}(0)\right]$$

Proof (Continued)

$$\mathbb{E}\left[d\mathbf{x}|\mathbf{x}(0)\right] = \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right] dt$$

$$\frac{d\mathbb{E}\left[\mathbf{x}(t)|\mathbf{x}(0)\right]}{dt} = \frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right]$$

Examples

$$\textbf{NCSN:} \quad \mathbf{f}(\mathbf{x}, t) = 0 \quad \Rightarrow \quad \boldsymbol{\mu} = \mathbf{x}(0)$$

$$\textbf{DDPM:} \quad \mathbf{f}(\mathbf{x}, t) = -\frac{1}{2}\beta(t)\mathbf{x}(t) \quad \Rightarrow \quad \frac{d\boldsymbol{\mu}}{dt} = -\frac{1}{2}\beta(t)\boldsymbol{\mu}$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}(t), t)|\mathbf{x}(0)\right]$$

Proof (Continued)

$$\mathbb{E}\left[d\mathbf{x}|\mathbf{x}(0)\right] = \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right] dt$$

$$\frac{d\mathbb{E}\left[\mathbf{x}(t)|\mathbf{x}(0)\right]}{dt} = \frac{d\boldsymbol{\mu}(\mathbf{x}(t), \mathbf{x}(0))}{dt} = \mathbb{E}\left[\mathbf{f}(\mathbf{x}, t)|\mathbf{x}(0)\right]$$

Examples

$$\textbf{NCSN:} \quad \mathbf{f}(\mathbf{x}, t) = 0 \quad \Rightarrow \quad \boldsymbol{\mu} = \mathbf{x}(0)$$

$$\textbf{DDPM:} \quad \mathbf{f}(\mathbf{x}, t) = -\frac{1}{2}\beta(t)\mathbf{x}(t) \quad \Rightarrow \quad \frac{d\boldsymbol{\mu}}{dt} = -\frac{1}{2}\beta(t)\boldsymbol{\mu}$$

$$\boldsymbol{\mu} = \mathbf{x}(0)\exp\left(-\frac{1}{2}\int_0^t \beta(s)ds\right)$$

Särkkä S., Solin A. Applied stochastic differential equations, 2019

Score-Based Generative Models Through SDEs

Training

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))}\mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))}\big\|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}(t),t)-\nabla_{\mathbf{x}(t)}\log q(\mathbf{x}(t)|\mathbf{x}(0))\big\|_2^2$$

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\Big(\boldsymbol{\mu}(\mathbf{x}(t),\mathbf{x}(0)),\boldsymbol{\Sigma}(\mathbf{x}(t),\mathbf{x}(0))\Big)$$

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Score-Based Generative Models Through SDEs

Training

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))}\mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))}\big\|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}(t),t)-\nabla_{\mathbf{x}(t)}\log q(\mathbf{x}(t)|\mathbf{x}(0))\big\|_2^2$$

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\Big(\boldsymbol{\mu}(\mathbf{x}(t),\mathbf{x}(0)),\boldsymbol{\Sigma}(\mathbf{x}(t),\mathbf{x}(0))\Big)$$

NCSN

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\mathbf{x}(0),\left[\sigma^2(t)-\sigma^2(0)\right]\cdot\mathbf{I}\right)$$

Score-Based Generative Models Through SDEs

Training

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))}\mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))}\big\|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}(t),t)-\nabla_{\mathbf{x}(t)}\log q(\mathbf{x}(t)|\mathbf{x}(0))\big\|_2^2$$

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\Big(\boldsymbol{\mu}(\mathbf{x}(t),\mathbf{x}(0)),\boldsymbol{\Sigma}(\mathbf{x}(t),\mathbf{x}(0))\Big)$$

NCSN

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\mathbf{x}(0), \left[\sigma^2(t) - \sigma^2(0)\right] \cdot \mathbf{I}\right)$$

DDPM

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\mathbf{x}(0)e^{-\frac{1}{2}\int_0^t \beta(s)ds}, \left(1 - e^{-\int_0^t \beta(s)ds}\right) \cdot \mathbf{I}\right)$$

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020

Score-Based Generative Models Through SDEs

Training

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))}\mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))}\big\|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}(t),t)-\nabla_{\mathbf{x}(t)}\log q(\mathbf{x}(t)|\mathbf{x}(0))\big\|_2^2$$

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\Big(\boldsymbol{\mu}(\mathbf{x}(t),\mathbf{x}(0)), \boldsymbol{\Sigma}(\mathbf{x}(t),\mathbf{x}(0))\Big)$$

NCSN

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\mathbf{x}(0), \left[\sigma^2(t) - \sigma^2(0)\right] \cdot \mathbf{I}\right)$$

DDPM

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\mathbf{x}(0)e^{-\frac{1}{2}\int_0^t \beta(s)ds}, \left(1 - e^{-\int_0^t \beta(s)ds}\right) \cdot \mathbf{I}\right)$$
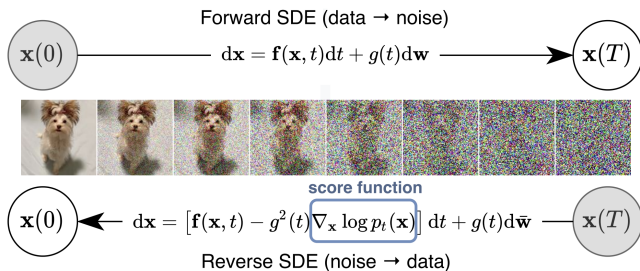
Here we omit the derivations of the variance.

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Score-Based Generative Models Through SDEs

Sampling

Solve the reverse SDE using numerical solvers (`SDESolve`).

Forward SDE (data → noise)

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

$\mathbf{x}(0)$ $\mathbf{x}(T)$

score function

$$d\mathbf{x} = \left[\mathbf{f}(\mathbf{x}, t) - g^2(t)\boxed{\nabla_{\mathbf{x}} \log p_t(\mathbf{x})}\right] dt + g(t)d\bar{\mathbf{w}}$$

Reverse SDE (noise → data)

Song Y., et al. Score-Based Generative Modeling through Stochastic Differential Equations, 2020
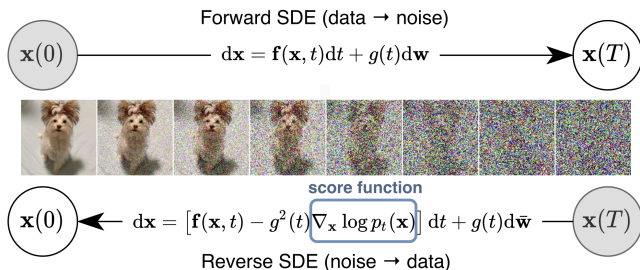
Score-Based Generative Models Through SDEs

Sampling

Solve the reverse SDE using numerical solvers (`SDESolve`).

Forward SDE (data → noise)

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

score function

$$d\mathbf{x} = \left[\mathbf{f}(\mathbf{x}, t) - g^2(t)\boxed{\nabla_\mathbf{x} \log p_t(\mathbf{x})}\right]dt + g(t)d\bar{\mathbf{w}}$$

Reverse SDE (noise → data)

▶ Discretizing the reverse SDE provides ancestral sampling.

▶ Discretizing the probability flow ODE yields deterministic sampling.

Song Y., et al. *Score-Based Generative Modeling through Stochastic Differential Equations*, 2020

Outline

Continuous-Time Normalizing Flows

Let's return to ODE dynamics $\mathbf{x}(t)$ in the interval $t \in [0, 1]$:

- $\mathbf{x}_0 \sim p_0(\mathbf{x}) = p(\mathbf{x})$, $\mathbf{x}_1 \sim p_1(\mathbf{x}) = p_{\text{data}}(\mathbf{x})$;
- $p(\mathbf{x})$ is a base distribution (e.g., $\mathcal{N}(0, \mathbf{I})$), and $p_{\text{data}}(\mathbf{x})$ is the true data distribution.

$$\frac{d\mathbf{x}}{dt} = \mathbf{f}(\mathbf{x}, t), \quad \text{with initial condition } \mathbf{x}(0) = \mathbf{x}_0.$$

Chen R. T. Q. et al. Neural Ordinary Differential Equations, 2018

Continuous-Time Normalizing Flows

Let's return to ODE dynamics $\mathbf{x}(t)$ in the interval $t \in [0, 1]$:

- $\mathbf{x}_0 \sim p_0(\mathbf{x}) = p(\mathbf{x})$, $\mathbf{x}_1 \sim p_1(\mathbf{x}) = p_{\text{data}}(\mathbf{x})$;
- $p(\mathbf{x})$ is a base distribution (e.g., $\mathcal{N}(0, \mathbf{I})$), and $p_{\text{data}}(\mathbf{x})$ is the true data distribution.

$$\frac{d\mathbf{x}}{dt} = \mathbf{f}(\mathbf{x}, t), \quad \text{with initial condition } \mathbf{x}(0) = \mathbf{x}_0.$$

KFP Theorem (Continuity Equation)

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, t) p_t(\mathbf{x})\right) \Leftrightarrow \frac{d \log p_t(\mathbf{x}(t))}{dt} = -\text{tr}\left(\frac{\partial \mathbf{f}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)}\right)$$

Chen R. T. Q. et al. Neural Ordinary Differential Equations, 2018

Continuous-Time Normalizing Flows

Let's return to ODE dynamics $\mathbf{x}(t)$ in the interval $t \in [0, 1]$:

- $\mathbf{x}_0 \sim p_0(\mathbf{x}) = p(\mathbf{x})$, $\mathbf{x}_1 \sim p_1(\mathbf{x}) = p_{\text{data}}(\mathbf{x})$;
- $p(\mathbf{x})$ is a base distribution (e.g., $\mathcal{N}(0, \mathbf{I})$), and $p_{\text{data}}(\mathbf{x})$ is the true data distribution.
$$\frac{d\mathbf{x}}{dt} = \mathbf{f}(\mathbf{x}, t), \quad \text{with initial condition } \mathbf{x}(0) = \mathbf{x}_0.$$

KFP Theorem (Continuity Equation)

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x})\right) \Leftrightarrow \frac{d \log p_t(\mathbf{x}(t))}{dt} = -\text{tr}\left(\frac{\partial \mathbf{f}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)}\right)$$

- It's hard to solve the continuity equation directly due to the trace term.
- There's a method (the adjoint method) that solves this equation directly, but it's unstable and unscalable.

Chen R. T. Q. et al. Neural Ordinary Differential Equations, 2018

Continuous-Time Normalizing Flows

KFP Theorem (Continuity Equation)

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, t) p_t(\mathbf{x})\right) \Leftrightarrow \frac{d \log p_t(\mathbf{x}(t))}{dt} = -\text{tr}\left(\frac{\partial \mathbf{f}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)}\right)$$

▶ Knowing the vector field $\mathbf{f}(\mathbf{x}, t)$, the KFP (or continuity) equation allows us to compute the density $p_t(\mathbf{x})$.

▶ Flow matching provides an alternative approach to Neural ODEs.

Lipman Y., et al. Flow Matching for Generative Modeling, 2022

Continuous-Time Normalizing Flows

KFP Theorem (Continuity Equation)

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, t) p_t(\mathbf{x})\right) \Leftrightarrow \frac{d \log p_t(\mathbf{x}(t))}{dt} = -\text{tr}\left(\frac{\partial \mathbf{f}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)}\right)$$

- ▶ Knowing the vector field $\mathbf{f}(\mathbf{x}, t)$, the KFP (or continuity) equation allows us to compute the density $p_t(\mathbf{x})$.
- ▶ Flow matching provides an alternative approach to Neural ODEs.

Flow Matching

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 \to \min_{\boldsymbol{\theta}}$$

Lipman Y., et al. Flow Matching for Generative Modeling, 2022

Continuous-Time Normalizing Flows

KFP Theorem (Continuity Equation)

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x})\right) \Leftrightarrow \frac{d\log p_t(\mathbf{x}(t))}{dt} = -\text{tr}\left(\frac{\partial \mathbf{f}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)}\right)$$

▶ Knowing the vector field $\mathbf{f}(\mathbf{x}, t)$, the KFP (or continuity) equation allows us to compute the density $p_t(\mathbf{x})$.

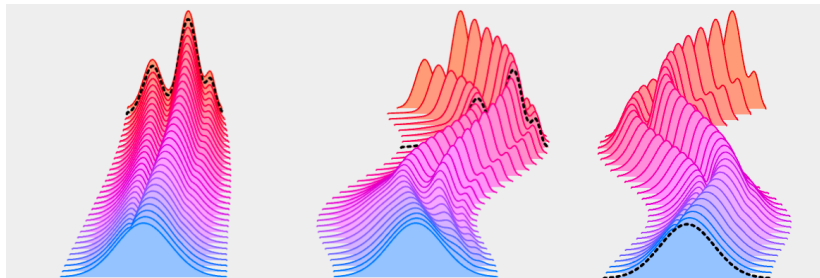▶ Flow matching provides an alternative approach to Neural ODEs.

Flow Matching

$$\mathbb{E}_{t \sim U[0,1]}\mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 \to \min_{\boldsymbol{\theta}}$$

▶ Approximate the true vector field $\mathbf{f}(\mathbf{x}, t)$ using $\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)$.

▶ Use $\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)$ for deterministic sampling from the ODE.

Lipman Y., et al. Flow Matching for Generative Modeling, 2022

Flow Matching

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 \to \min_{\boldsymbol{\theta}}$$

▶ There are infinitely many possible $\mathbf{f}(\mathbf{x}, t)$ between $p_{\text{data}}(\mathbf{x})$ and $p(\mathbf{x})$.

▶ The true vector field $\mathbf{f}(\mathbf{x}, t)$ is **unknown**.

▶ We need to select the "best" $\mathbf{f}(\mathbf{x}, t)$ and make the objective tractable.

image credit: A Visual Dive into Conditional Flow Matching

Outline

Flow Matching

Latent Variable Model

Let's introduce the latent variable $\mathbf{z}$:

$$p_t(\mathbf{x}) = \int p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

Here, $p_t(\mathbf{x}|\mathbf{z})$ is a **conditional probability path**.

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport, 2023*

Flow Matching

Latent Variable Model

Let's introduce the latent variable $\mathbf{z}$:

$$p_t(\mathbf{x}) = \int p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

Here, $p_t(\mathbf{x}|\mathbf{z})$ is a **conditional probability path**. The conditional probability path $p_t(\mathbf{x}|\mathbf{z})$ satisfies the KFP theorem:

$$\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})\right),$$

where $\mathbf{f}(\mathbf{x}, \mathbf{z}, t)$ is a **conditional vector field**:

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport*, 2023

Flow Matching

Latent Variable Model

Let's introduce the latent variable $\mathbf{z}$:

$$p_t(\mathbf{x}) = \int p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

Here, $p_t(\mathbf{x}|\mathbf{z})$ is a **conditional probability path**. The conditional probability path $p_t(\mathbf{x}|\mathbf{z})$ satisfies the KFP theorem:

$$\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})\right),$$

where $\mathbf{f}(\mathbf{x}, \mathbf{z}, t)$ is a **conditional vector field**:

$$\frac{d\mathbf{x}}{dt} = \mathbf{f}(\mathbf{x}, t) \quad \Rightarrow \quad \frac{d\mathbf{x}}{dt} = \mathbf{f}(\mathbf{x}, \mathbf{z}, t)$$

What's the relationship between $\mathbf{f}(\mathbf{x}, t)$ and $\mathbf{f}(\mathbf{x}, \mathbf{z}, t)$?

Tong A., et al. Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport, 2023

Flow Matching

$$\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t} = -\mathrm{div}\left(\mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})\right),$$

Theorem

The following vector field generates the probability path $p_t(\mathbf{x})$:

$$\mathbf{f}(\mathbf{x}, t) = \mathbb{E}_{p_t(\mathbf{z}|\mathbf{x})}\mathbf{f}(\mathbf{x}, \mathbf{z}, t) = \int \mathbf{f}(\mathbf{x}, \mathbf{z}, t)\frac{p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p_t(\mathbf{x})}d\mathbf{z}$$

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport*, 2023

Flow Matching

$$\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})\right),$$

Theorem
The following vector field generates the probability path $p_t(\mathbf{x})$:

$$\mathbf{f}(\mathbf{x}, t) = \mathbb{E}_{p_t(\mathbf{z}|\mathbf{x})}\mathbf{f}(\mathbf{x}, \mathbf{z}, t) = \int \mathbf{f}(\mathbf{x}, \mathbf{z}, t)\frac{p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p_t(\mathbf{x})}d\mathbf{z}$$

Proof

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = \frac{\partial}{\partial t}\int p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

Tong A., et al. Improving and Generalizing Flow-Based Generative Models with
Minibatch Optimal Transport, 2023

Flow Matching

$$\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})\right),$$

Theorem

The following vector field generates the probability path $p_t(\mathbf{x})$:

$$\mathbf{f}(\mathbf{x}, t) = \mathbb{E}_{p_t(\mathbf{z}|\mathbf{x})}\mathbf{f}(\mathbf{x}, \mathbf{z}, t) = \int \mathbf{f}(\mathbf{x}, \mathbf{z}, t)\frac{p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p_t(\mathbf{x})}d\mathbf{z}$$

Proof

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = \frac{\partial}{\partial t}\int p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z} = \int \left(\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t}\right)p(\mathbf{z})d\mathbf{z}$$

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport*, 2023

Flow Matching

$$\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})\right),$$

Theorem
The following vector field generates the probability path $p_t(\mathbf{x})$:

$$\mathbf{f}(\mathbf{x}, t) = \mathbb{E}_{p_t(\mathbf{z}|\mathbf{x})}\mathbf{f}(\mathbf{x}, \mathbf{z}, t) = \int \mathbf{f}(\mathbf{x}, \mathbf{z}, t)\frac{p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p_t(\mathbf{x})}d\mathbf{z}$$

Proof

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = \frac{\partial}{\partial t}\int p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z} = \int \left(\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t}\right) p(\mathbf{z})d\mathbf{z} =$$

$$= \int \left(-\text{div}\left(\mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})\right)\right) p(\mathbf{z})d\mathbf{z}$$

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport*, 2023

Flow Matching

$$\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})\right),$$

Theorem

The following vector field generates the probability path $p_t(\mathbf{x})$:

$$\mathbf{f}(\mathbf{x}, t) = \mathbb{E}_{p_t(\mathbf{z}|\mathbf{x})}\mathbf{f}(\mathbf{x}, \mathbf{z}, t) = \int \mathbf{f}(\mathbf{x}, \mathbf{z}, t)\frac{p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p_t(\mathbf{x})}d\mathbf{z}$$

Proof

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = \frac{\partial}{\partial t}\int p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z} = \int \left(\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t}\right)p(\mathbf{z})d\mathbf{z} =$$

$$= \int \left(-\text{div}\left(\mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})\right)\right)p(\mathbf{z})d\mathbf{z} =$$

$$= -\text{div}\left(\int \mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}\right)$$

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport*, 2023

Flow Matching

$$\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t} = -\text{div}\left(\mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})\right),$$

Theorem
The following vector field generates the probability path $p_t(\mathbf{x})$:

$$\mathbf{f}(\mathbf{x}, t) = \mathbb{E}_{p_t(\mathbf{z}|\mathbf{x})}\mathbf{f}(\mathbf{x}, \mathbf{z}, t) = \int \mathbf{f}(\mathbf{x}, \mathbf{z}, t)\frac{p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p_t(\mathbf{x})}d\mathbf{z}$$

Proof

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = \frac{\partial}{\partial t}\int p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z} = \int \left(\frac{\partial p_t(\mathbf{x}|\mathbf{z})}{\partial t}\right)p(\mathbf{z})d\mathbf{z} =$$

$$= \int \left(-\text{div}\left(\mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})\right)\right)p(\mathbf{z})d\mathbf{z} =$$

$$= -\text{div}\left(\int \mathbf{f}(\mathbf{x}, \mathbf{z}, t)p_t(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}\right) = -\text{div}\left(\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x})\right)$$

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport*, 2023

Flow Matching

Flow Matching (FM)

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 \to \min_{\boldsymbol{\theta}}$$

Conditional Flow Matching (CFM)

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x}|\mathbf{z})} \|\mathbf{f}(\mathbf{x}, \mathbf{z}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 \to \min_{\boldsymbol{\theta}}$$

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport*, 2023

Flow Matching

Flow Matching (FM)

$$\mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x})}\left\|\mathbf{f}(\mathbf{x},t)-\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x},t)\right\|^2 \to \min_{\boldsymbol{\theta}}$$

Conditional Flow Matching (CFM)

$$\mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{z}\sim p(\mathbf{z})}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x}|\mathbf{z})}\left\|\mathbf{f}(\mathbf{x},\mathbf{z},t)-\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x},t)\right\|^2 \to \min_{\boldsymbol{\theta}}$$

Theorem
If $\text{supp}(p_t(\mathbf{x})) = \mathbb{R}^m$, then the optimal value of the FM objective equals the optimal value of the CFM objective.

Tong A., et al. Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport, 2023

Flow Matching

Flow Matching (FM)

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 \to \min_{\boldsymbol{\theta}}$$

Conditional Flow Matching (CFM)

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x}|\mathbf{z})} \|\mathbf{f}(\mathbf{x}, \mathbf{z}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 \to \min_{\boldsymbol{\theta}}$$

Theorem
If $\mathrm{supp}(p_t(\mathbf{x})) = \mathbb{R}^m$, then the optimal value of the FM objective equals the optimal value of the CFM objective.

Proof
This can be proved in a similar way as in the denoising score matching theorem.

Tong A., et al. Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport, 2023

Conditional Flow Matching

Theorem

$$\arg\min_{\boldsymbol{\theta}} \mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 =$$

$$= \arg\min_{\boldsymbol{\theta}} \mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{z}\sim p(\mathbf{z})}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x}|\mathbf{z})} \|\mathbf{f}(\mathbf{x}, \mathbf{z}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2$$

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport*, 2023

Conditional Flow Matching

Theorem

$$\arg\min_{\boldsymbol{\theta}} \mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x})} \left\| \mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t) \right\|^2 =$$

$$= \arg\min_{\boldsymbol{\theta}} \mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{z}\sim p(\mathbf{z})}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x}|\mathbf{z})} \left\| \mathbf{f}(\mathbf{x}, \mathbf{z}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t) \right\|^2$$

Proof

$$\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x})} \left\| \mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t) \right\|^2 =$$

$$= \mathbb{E}_{\mathbf{z}\sim p(\mathbf{z})}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x}|\mathbf{z})} \left[\left\| \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t) \right\|^2 - 2\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x}, t)\mathbf{f}(\mathbf{x}, t) \right] + \mathrm{const}(\boldsymbol{\theta})$$

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport*, 2023

Conditional Flow Matching

Theorem

$$\arg\min_{\boldsymbol{\theta}} \mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x},t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x},t)\|^2 =$$

$$= \arg\min_{\boldsymbol{\theta}} \mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{z}\sim p(\mathbf{z})}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x}|\mathbf{z})} \|\mathbf{f}(\mathbf{x},\mathbf{z},t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x},t)\|^2$$

Proof

$$\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x},t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x},t)\|^2 =$$

$$= \mathbb{E}_{\mathbf{z}\sim p(\mathbf{z})}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x}|\mathbf{z})} \left[\|\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x},t)\|^2 - 2\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x},t)\mathbf{f}(\mathbf{x},t) \right] + \mathrm{const}(\boldsymbol{\theta})$$

$$\mathbb{E}_{p_t(\mathbf{x})} \left[\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x},t)\mathbf{f}(\mathbf{x},t) \right] = \int p_t(\mathbf{x}) \left[\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x},t) \left(\int p_t(\mathbf{z}|\mathbf{x})\mathbf{f}(\mathbf{x},\mathbf{z},t)d\mathbf{z} \right) \right] d\mathbf{x}$$

Tong A., et al. Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport, 2023

Conditional Flow Matching

Theorem

$$\arg\min_{\boldsymbol{\theta}} \mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 =$$

$$= \arg\min_{\boldsymbol{\theta}} \mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{z}\sim p(\mathbf{z})}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x}|\mathbf{z})} \|\mathbf{f}(\mathbf{x}, \mathbf{z}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2$$

Proof

$$\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 =$$

$$= \mathbb{E}_{\mathbf{z}\sim p(\mathbf{z})}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x}|\mathbf{z})} \left[\|\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 - 2\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x}, t)\mathbf{f}(\mathbf{x}, t) \right] + \text{const}(\boldsymbol{\theta})$$

$$\mathbb{E}_{p_t(\mathbf{x})} \left[\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x}, t)\mathbf{f}(\mathbf{x}, t) \right] = \int p_t(\mathbf{x}) \left[\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x}, t) \left(\int p_t(\mathbf{z}|\mathbf{x})\mathbf{f}(\mathbf{x}, \mathbf{z}, t)d\mathbf{z} \right) \right] d\mathbf{x} =$$

$$= \int \int p_t(\mathbf{x})p_t(\mathbf{z}|\mathbf{x}) \left[\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x}, t)\mathbf{f}(\mathbf{x}, \mathbf{z}, t) \right] d\mathbf{z}d\mathbf{x}$$

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport*, 2023
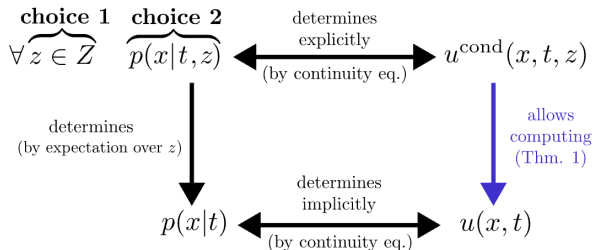
Conditional Flow Matching

Theorem

$$\arg\min_{\boldsymbol{\theta}} \mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 =$$

$$= \arg\min_{\boldsymbol{\theta}} \mathbb{E}_{t\sim U[0,1]}\mathbb{E}_{\mathbf{z}\sim p(\mathbf{z})}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x}|\mathbf{z})} \|\mathbf{f}(\mathbf{x}, \mathbf{z}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2$$
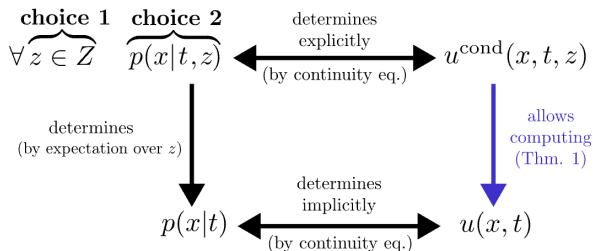
Proof

$$\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x})} \|\mathbf{f}(\mathbf{x}, t) - \mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 =$$

$$= \mathbb{E}_{\mathbf{z}\sim p(\mathbf{z})}\mathbb{E}_{\mathbf{x}\sim p_t(\mathbf{x}|\mathbf{z})} \left[\|\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}, t)\|^2 - 2\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x}, t)\mathbf{f}(\mathbf{x}, t) \right] + \mathrm{const}(\boldsymbol{\theta})$$

$$\mathbb{E}_{p_t(\mathbf{x})} \left[\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x}, t)\mathbf{f}(\mathbf{x}, t) \right] = \int p_t(\mathbf{x}) \left[\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x}, t) \left(\int p_t(\mathbf{z}|\mathbf{x})\mathbf{f}(\mathbf{x}, \mathbf{z}, t)d\mathbf{z} \right) \right] d\mathbf{x} =$$

$$= \int \int p_t(\mathbf{x})p_t(\mathbf{z}|\mathbf{x}) \left[\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x}, t)\mathbf{f}(\mathbf{x}, \mathbf{z}, t) \right] d\mathbf{z}d\mathbf{x} =$$

$$= \mathbb{E}_{p(\mathbf{z})}\mathbb{E}_{p_t(\mathbf{x}|\mathbf{z})} \left[\mathbf{f}_{\boldsymbol{\theta}}^T(\mathbf{x}, t)\mathbf{f}(\mathbf{x}, \mathbf{z}, t) \right]$$

Tong A., et al. *Improving and Generalizing Flow-Based Generative Models with Minibatch Optimal Transport*, 2023

Conditional Flow Matching

Conditional Flow Matching

- We don't want to directly model $p_t(\mathbf{x})$, since it's complex.
- We've shown it's possible to solve the CFM task instead of the FM task.
- Let's choose a convenient conditioning latent variable $\mathbf{z}$.
- We'll parametrize $p_t(\mathbf{x}|\mathbf{z})$ instead of $p_t(\mathbf{x})$. It should satisfy the following constraints:

$$p(\mathbf{x}) = \mathcal{N}(0, \mathbf{I}) = \mathbb{E}_{p(\mathbf{z})} p_0(\mathbf{x}|\mathbf{z}); \quad p_{\text{data}}(\mathbf{x}) = \mathbb{E}_{p(\mathbf{z})} p_1(\mathbf{x}|\mathbf{z}).$$

image credit: A Visual Dive into Conditional Flow Matching

Summary

▶ Score matching (NCSN) and diffusion models (DDPM) are discretizations of SDEs (variance exploding and variance preserving).

▶ It's possible to train continuous-in-time score-based generative models using forward and reverse SDEs.

▶ Discretizing the reverse SDE yields ancestral sampling of the DDPM.

▶ Flow matching suggests fitting the vector field directly.

▶ Conditional flow matching introduces the latent variable z, reformulating the initial task in terms of conditional dynamics.